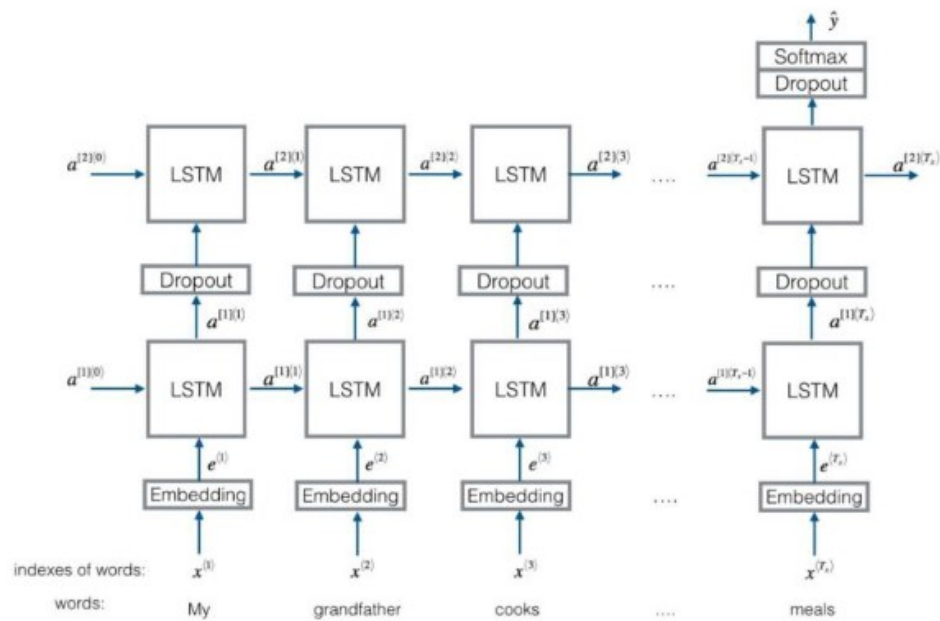


HW5 Report

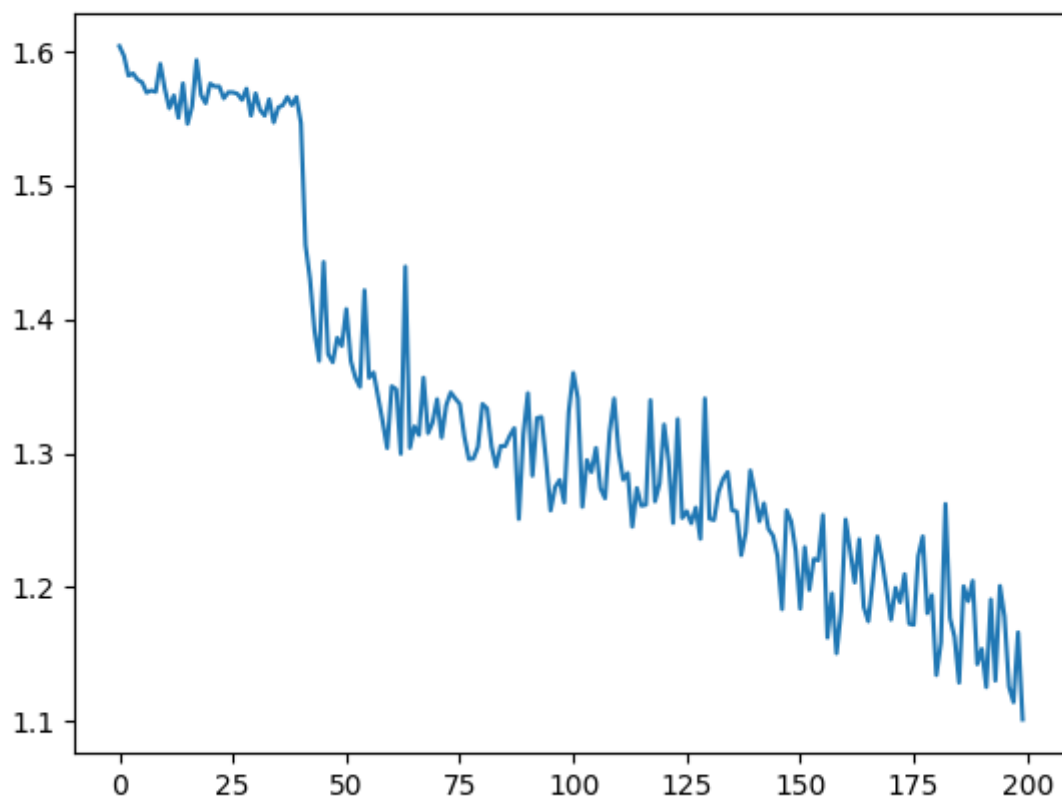
Architecture



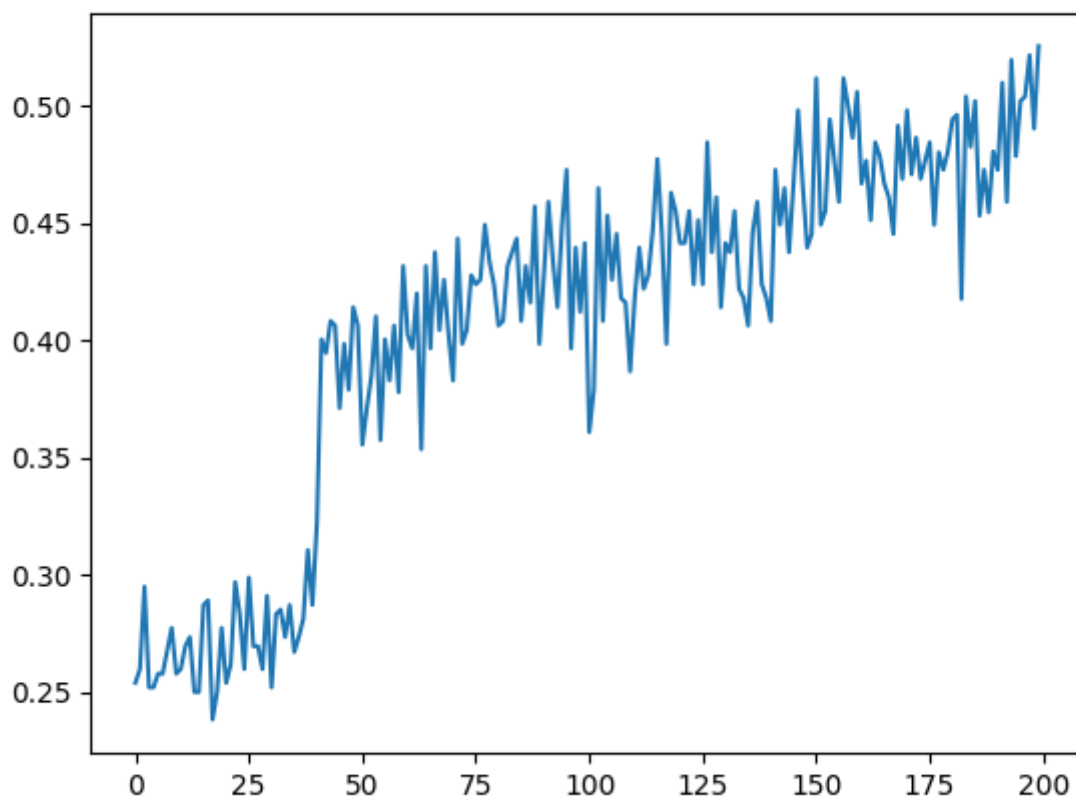
After embedding, data passes through 2 LSTM layers (bidirectional) and the get out of the last LSTM forwarding output into a FCLayer.

Records

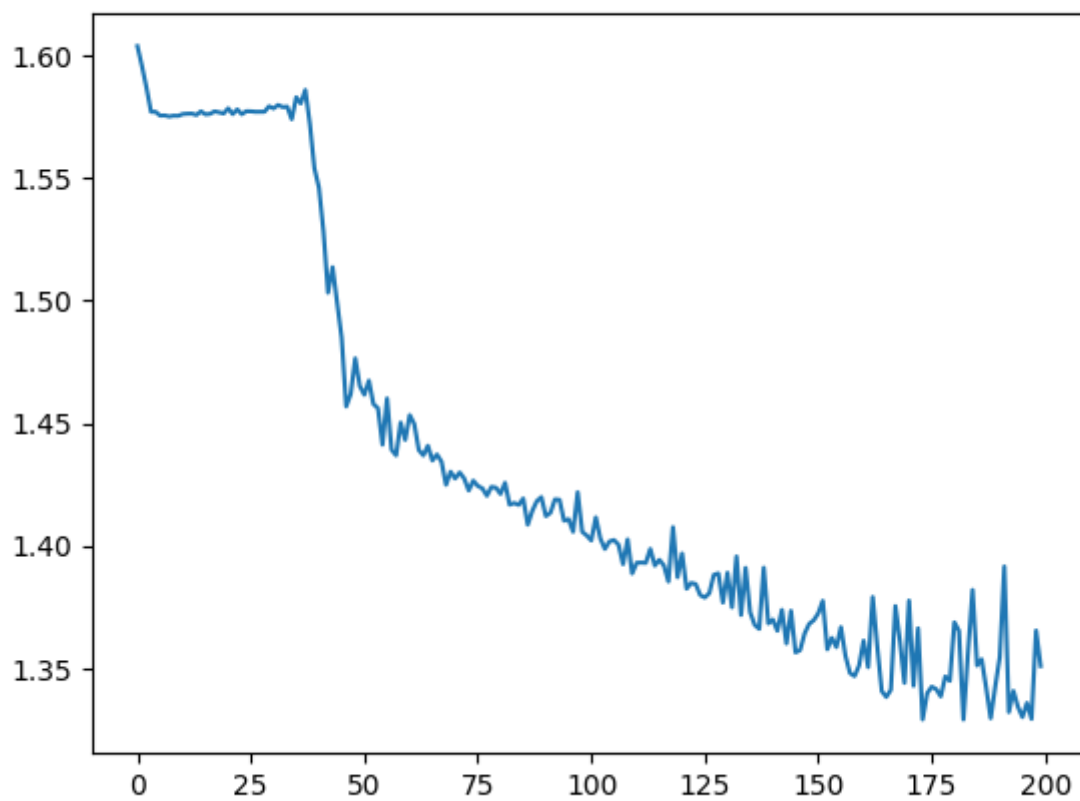
Training Performance



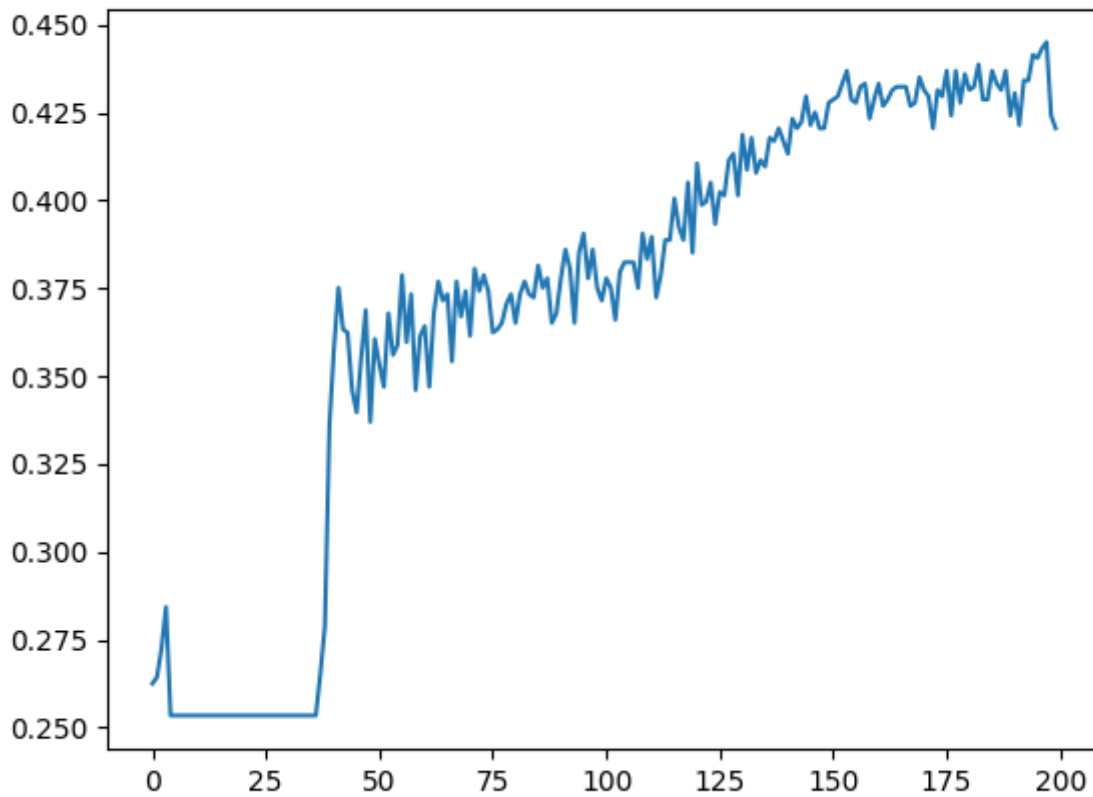
training loss



training accuracy



validation loss



validation accuracy

LAST VALIDATION LOSS

1.351

LAST VALIDATION ACCURACY

42.05%

Test result

TEST LOSS

1.318

TEST ACCURACY

43.35%

Hyperparameters

```
# hyperparameters
learning_rate = 0.0001
num_epochs = 200
batch_size = 512
lbd = 0.0001
dropout = 0.5
n_layers = 2

EMBEDDING_DIM = 300
HIDDEN_DIM = 128
OUTPUT_DIM = 5
BIDIRECTIONAL = True
```

In the plots, though it's generally converging, the bouncing is seriously violently (even less smooth in the beginning). I enlarged `batch_size` and decreased learning rate, and used Adam as optimizer to make it go more smoothly. But still, as the data set is limited, that's hardly to avoid.